

Things to do

- Problem Statement
- Data set (min 7000 rows)
- Handling missing data
- Handling feature scaling
- Best possible methods (Algorithms) for our dataset
- Visualization

Dataset Overview

Nearest Earth Objects (NEO)

Link - <https://www.kaggle.com/datasets/ivansher/nasa-nearest-earth-objects-1910-2024/data>

This dataset contains **338,199** records of objects that come close to Earth, known as **NEOs (Near-Earth Objects)**, recorded between **1910 and 2024**. NASA monitors these objects to assess their potential risk to the planet.

Key Features

- **neo_id**: Unique identifier for each NEO.
- **name**: Designation or name of the object.
- **absolute_magnitude**: The intrinsic brightness of the object.
- **estimated_diameter_min/max**: Estimated size range of the object.
- **relative_velocity**: The object's speed relative to Earth (in km/h).
- **miss_distance**: How close the object will come to Earth (in km).
- **is_hazardous**: A boolean field indicating if NASA classifies the object as potentially dangerous.

Problem Statement

Predicting Potentially Hazardous NEOs The goal is to build a machine learning model to predict whether a **Near-Earth Object (NEO)** is classified as "**hazardous**" by NASA.

Objectives

1. **Preprocessing**: Clean the dataset and handle missing values.
2. **Feature Selection**: Used key features such as:
 - Absolute magnitude
 - Estimated diameter
 - Relative velocity
 - Miss distance
3. **Oversampling**: Address class imbalance by applying **SMOTE** (Synthetic Minority Oversampling Technique) to balance the target class (**is_hazardous**).
4. **Modeling**: Train a model to predict the **is_hazardous** label (True/False).

5. **Evaluation:** Use metrics such as **accuracy**, **precision**, **recall**, and **ROC-AUC score** to evaluate the model.

```
import pandas as pd
df = pd.read_csv('../data/nearest-earth-objects.csv')

df.head()
```

	neo_id		name	absolute_magnitude
	estimated_diameter_min \			
0	2162117	162117	(1998 SD15)	19.14
	0.394962			
1	2349507	349507	(2008 QY)	18.50
	0.530341			
2	2455415	455415	(2003 GA)	21.45
	0.136319			
3	3132126		(2002 PB)	20.63
	0.198863			
4	3557844		(2011 DW)	22.70
	0.076658			

	estimated_diameter_max	orbiting_body	relative_velocity	
	miss_distance \			
0		0.883161	Earth	71745.401048
	5.814362e+07			
1		1.185878	Earth	109949.757148
	5.580105e+07			
2		0.304818	Earth	24865.506798
	6.720689e+07			
3		0.444672	Earth	78890.076805
	3.039644e+07			
4		0.171412	Earth	56036.519484
	6.311863e+07			

	is_hazardous
0	False
1	True
2	False
3	False
4	False

Dataset Preprocessing

```
# Remove rows with missing values
df_cleaned = df.dropna()

missing_values = df_cleaned.isnull()

# To get a summary of missing values in each column:
missing_values_summary = missing_values.sum()
```

missing_values_summary

neo_id	0
name	0
absolute_magnitude	0
estimated_diameter_min	0
estimated_diameter_max	0
orbiting_body	0
relative_velocity	0
miss_distance	0
is_hazardous	0
dtype: int64	

Dropping unnecessary columns

```
# Drop 'neo_id', 'name', and 'orbiting_body' since they are not useful for prediction
df_cleaned = df_cleaned.drop(columns=['neo_id', 'name', 'orbiting_body'])
```

Feature Scaling using MinMaxScaler

```
from sklearn.preprocessing import MinMaxScaler

# Select the features to scale
features_to_scale = ['absolute_magnitude', 'estimated_diameter_min', 'estimated_diameter_max', 'relative_velocity', 'miss_distance']

# Initialize the scaler
scaler = MinMaxScaler()

# Apply scaling
df_cleaned[features_to_scale] = scaler.fit_transform(df_cleaned[features_to_scale])
```

The dataset contains 9 columns and 338171 rows of data

df_cleaned.shape

(338171, 6)

Heavily Skewed Data Set

```
target_column = df_cleaned['is_hazardous']

# Count the occurrences of True and False
true_count = target_column.value_counts()[1]
```

```

false_count = target_column.value_counts()[0]

true_count, false_count

C:\Users\amrit\AppData\Local\Temp\ipykernel_21944\1876494894.py:4:
FutureWarning: Series.__getitem__ treating keys as positions is
deprecated. In a future version, integer keys will always be treated
as labels (consistent with DataFrame behavior). To access a value by
position, use `ser.iloc[pos]`
    true_count = target_column.value_counts()[1]
C:\Users\amrit\AppData\Local\Temp\ipykernel_21944\1876494894.py:5:
FutureWarning: Series.__getitem__ treating keys as positions is
deprecated. In a future version, integer keys will always be treated
as labels (consistent with DataFrame behavior). To access a value by
position, use `ser.iloc[pos]`
    false_count = target_column.value_counts()[0]


(43162, 295009)

import matplotlib.pyplot as plt

# Create a bar plot to visualize the distribution of True and False
labels = ['False (Not Hazardous)', 'True (Hazardous)']
counts = [false_count, true_count]

plt.bar(labels, counts, color=['orange', 'green'])
plt.title('Distribution of Hazardous vs Non-Hazardous Objects')
plt.ylabel('Count')
plt.show()

```



Resampling (Oversampling) with SMOTE

```
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split

# Separate features and target
X = df_cleaned.drop(columns=['is_hazardous'])
y = df_cleaned['is_hazardous']

# Split the dataset into training and test sets before applying SMOTE
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Apply SMOTE only to the training set to avoid data leakage
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train,
y_train)

# Check the class distribution after SMOTE
print(y_train_resampled.value_counts())
```

```
is_hazardous
False      206396
True       206396
Name: count, dtype: int64
```

Training

Random Forest

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix,
roc_auc_score, accuracy_score

# Initialize the Random Forest Classifier
rf_model = RandomForestClassifier(random_state=42)

# Train the model on the resampled training set
rf_model.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)

# Evaluate the model
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, rf_model.predict_proba(X_test)[:, 1])

# Print the evaluation results
print("Confusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
print("Accuracy:", accuracy)
print("ROC AUC Score:", roc_auc)
```

Confusion Matrix:

```
[[82920  5693]
 [ 3853  8986]]
```

Classification Report:

	precision	recall	f1-score	support
False	0.96	0.94	0.95	88613
True	0.61	0.70	0.65	12839
accuracy			0.91	101452
macro avg	0.78	0.82	0.80	101452
weighted avg	0.91	0.91	0.91	101452

Accuracy: 0.9059062413752317

ROC AUC Score: 0.943386150661994

KNN

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix,
```

```

accuracy_score

# Initialize the KNN Classifier
knn_model = KNeighborsClassifier(n_neighbors=5)

# Train the model on the resampled training set
knn_model.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test set
y_pred_knn = knn_model.predict(X_test)

# Evaluate the model
conf_matrix = confusion_matrix(y_test, y_pred_knn)
class_report = classification_report(y_test, y_pred_knn)
accuracy = accuracy_score(y_test, y_pred_knn)

# Print the evaluation results
print("Confusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
print("Accuracy:", accuracy)

Confusion Matrix:
[[72904 15709]
 [ 3318  9521]]

Classification Report:

```

	precision	recall	f1-score	support
False	0.96	0.82	0.88	88613
True	0.38	0.74	0.50	12839
accuracy			0.81	101452
macro avg	0.67	0.78	0.69	101452
weighted avg	0.88	0.81	0.84	101452

```

Accuracy: 0.8124531798288845

```

Logistic Regression

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score, roc_auc_score

# Initialize the Logistic Regression model
logreg_model = LogisticRegression(random_state=42)

# Train the model on the resampled training set
logreg_model.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test set

```

```

y_pred_logreg = logreg_model.predict(X_test)

# Evaluate the model
conf_matrix_logreg = confusion_matrix(y_test, y_pred_logreg)
class_report_logreg = classification_report(y_test, y_pred_logreg)
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
roc_auc_logreg = roc_auc_score(y_test,
logreg_model.predict_proba(X_test)[: , 1])

# Print the evaluation results
print("Confusion Matrix:\n", conf_matrix_logreg)
print("\nClassification Report:\n", class_report_logreg)
print("Accuracy:", accuracy_logreg)
print("ROC AUC Score:", roc_auc_logreg)

Confusion Matrix:
[[63364 25249]
 [ 1542 11297]]

Classification Report:

```

	precision	recall	f1-score	support
False	0.98	0.72	0.83	88613
True	0.31	0.88	0.46	12839
accuracy			0.74	101452
macro avg	0.64	0.80	0.64	101452
weighted avg	0.89	0.74	0.78	101452

```

Accuracy: 0.73592437803099
ROC AUC Score: 0.8356701345776563

```

Decision Tree

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score, roc_auc_score

# Initialize the Decision Tree Classifier
dt_model = DecisionTreeClassifier(random_state=42)

# Train the model on the resampled training set
dt_model.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test set
y_pred_dt = dt_model.predict(X_test)

# Evaluate the model
conf_matrix_dt = confusion_matrix(y_test, y_pred_dt)
class_report_dt = classification_report(y_test, y_pred_dt)

```



```
accuracy_dt = accuracy_score(y_test, y_pred_dt)
roc_auc_dt = roc_auc_score(y_test, dt_model.predict_proba(X_test)[: ,
1])
```

Print the evaluation results

```
print("Confusion Matrix:\n", conf_matrix_dt)
print("\nClassification Report:\n", class_report_dt)
print("Accuracy:", accuracy_dt)
print("ROC AUC Score:", roc_auc_dt)
```

Confusion Matrix:

```
[[80745  7868]
 [ 4453  8386]]
```

Classification Report:

	precision	recall	f1-score	support
False	0.95	0.91	0.93	88613
True	0.52	0.65	0.58	12839
accuracy			0.88	101452
macro avg	0.73	0.78	0.75	101452
weighted avg	0.89	0.88	0.88	101452

Accuracy: 0.878553404565706

ROC AUC Score: 0.782187775329878

Visualization

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc, confusion_matrix

# Store accuracy and ROC AUC scores in a dictionary
model_performance = {
    'Random Forest': {'accuracy': 0.9059, 'roc_auc': 0.9434},
    'KNN': {'accuracy': 0.8125, 'roc_auc': 0.7868},
    'Logistic Regression': {'accuracy': 0.7359, 'roc_auc': 0.8357},
    'Decision Tree': {'accuracy': 0.8785, 'roc_auc': 0.7821},
}

# Extract model names and their corresponding accuracy and ROC-AUC
scores
model_names = list(model_performance.keys())
accuracies = [model_performance[model]['accuracy'] for model in
model_names]
roc_aucs = [model_performance[model]['roc_auc'] for model in
model_names]

# Plot the accuracies
```

```

plt.figure(figsize=(10, 5))
sns.barplot(x=model_names, y=accuracies, palette='viridis')
plt.title('Model Accuracy Comparison')
plt.ylabel('Accuracy')
plt.ylim(0, 1)
plt.show()

# Plot the ROC AUC scores
plt.figure(figsize=(10, 5))
sns.barplot(x=model_names, y=roc_aucs, palette='magma')
plt.title('Model ROC AUC Score Comparison')
plt.ylabel('ROC AUC Score')
plt.ylim(0, 1)
plt.show()

# Plot ROC Curves for all models (with y_test and predicted
probabilities for each model)
def plot_roc_curves(models, X_test, y_test):
    plt.figure(figsize=(10, 8))

    for name, model in models.items():
        # Predict probabilities
        if hasattr(model, "predict_proba"):
            y_proba = model.predict_proba(X_test)[:, 1] # Probability
            estimates for the positive class
        else:
            y_proba = model.decision_function(X_test) # For models
            like SVM that don't use predict_proba
        fpr, tpr, _ = roc_curve(y_test, y_proba)
        roc_auc = auc(fpr, tpr)
        plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

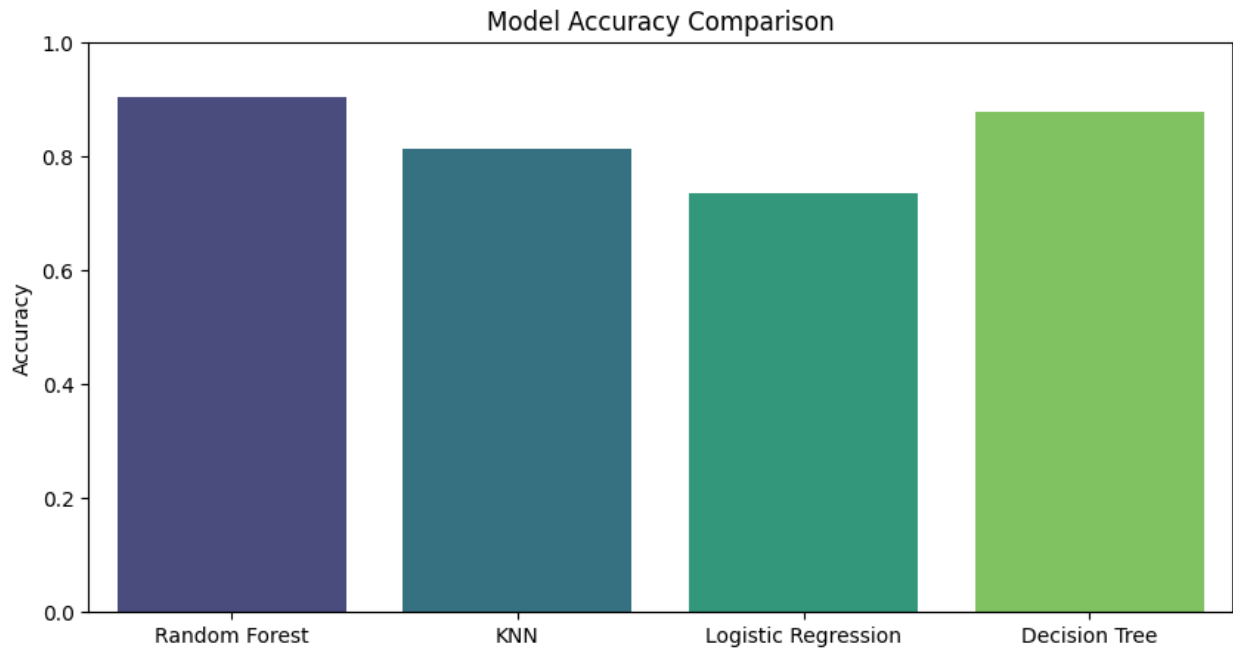
    plt.plot([0, 1], [0, 1], 'k--') # Diagonal line for random
    guessing
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curves Comparison')
    plt.legend(loc="lower right")
    plt.show()

```

C:\Users\amrit\AppData\Local\Temp\ipykernel_23972\3653782360.py:20:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

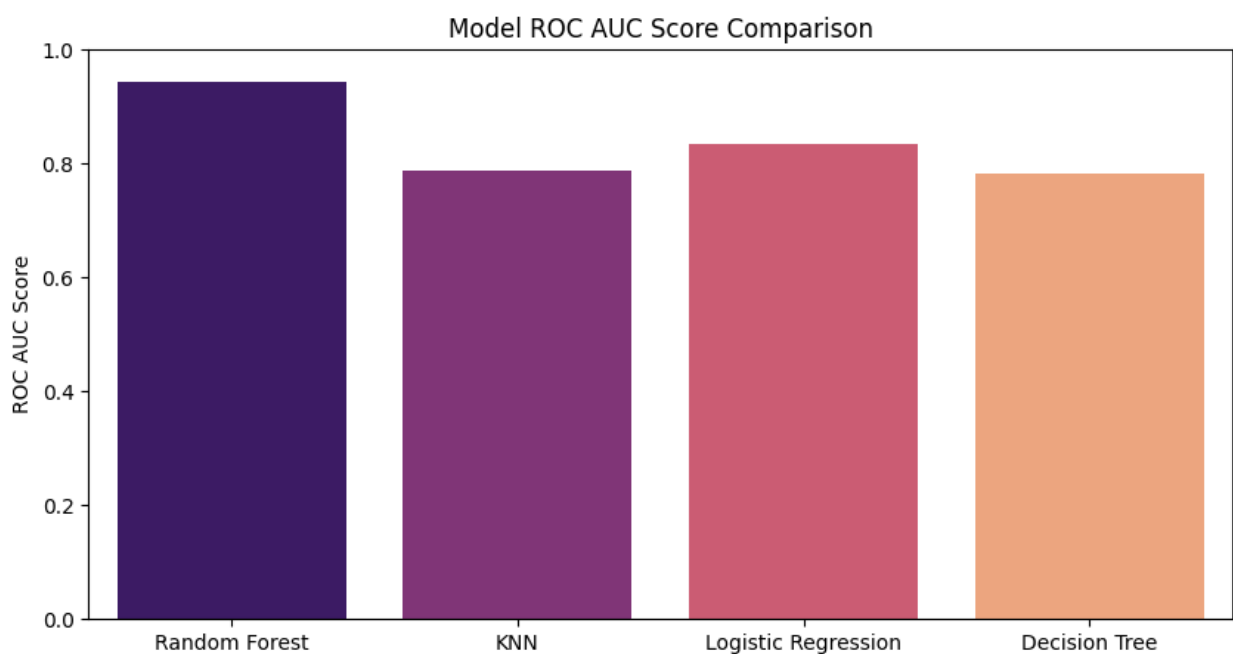
```
sns.barplot(x=model_names, y=accuracies, palette='viridis')
```



```
C:\Users\amrit\AppData\Local\Temp\ipykernel_23972\3653782360.py:28:  
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be  
removed in v0.14.0. Assign the `x` variable to `hue` and set  
`legend=False` for the same effect.
```

```
sns.barplot(x=model_names, y=roc_aucs, palette='magma')
```



```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import ConfusionMatrixDisplay

# Define the confusion matrices
confusion_matrices = {
    'Random Forest': [[82920, 5693], [3853, 8986]],
    'KNN': [[72904, 15709], [3318, 9521]],
    'Logistic Regression': [[63364, 25249], [1542, 11297]],
    'Decision Tree': [[80745, 7868], [4453, 8386]],
}

# Plotting function for confusion matrices
def plot_confusion_matrices(confusion_matrices):
    plt.figure(figsize=(12, 10))
    for i, (model, cm) in enumerate(confusion_matrices.items()):
        plt.subplot(2, 2, i + 1)
        sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
                    xticklabels=['False', 'True'],
                    yticklabels=['False', 'True'])
        plt.title(f'{model} Confusion Matrix')
        plt.xlabel('Predicted Label')
        plt.ylabel('True Label')

    plt.tight_layout()
    plt.show()

# Call the function to plot
plot_confusion_matrices(confusion_matrices)

```

